

# Context-dependent phone models


Peter Bell

Automatic Speech Recognition  
ASR Lecture 6  
30 January 2019

## Phone models

- Modelling phones with HMMs
- The need to model phonetic context
- Biphone and triphone models
- Parameter sharing – sharing parameters across different contexts
- Choosing which states to share – phonetic decision trees

# Recap: Continuous Density HMM




Probabilistic finite state automaton


Parameters  $\lambda$ :

- Transition probabilities:  $a_{kj} = P(q_{t+1} = j | q_t = k)$
- Output probability density function:  $b_j(x) = p(x | q = j)$

# Modelling words with HMM phone models




# Modelling words with HMM phone models



## Another example


Single emitting state per phone



“six”


# Another example

Three emitting states per phone




# Another example

Three emitting states per phone





# Word sequence model




“six quid”

- **Context** The acoustic phonetic context of a speech unit has an effect on its acoustic realization
- Consider /n/ in **ten** and **tenth**
  - alveolar in **ten**
  - dental in **tenth**
- **Coarticulation** the place of articulation for one speech sound depends on a neighbouring speech sound.
- Eg. /n/ in **ten pounds** vs. **ten kilos**
- More problematic in natural spontaneous speech


- **Context** The acoustic phonetic context of a speech unit has an effect on its acoustic realization
- Consider /n/ in **ten** and **tenth**
  - alveolar in **ten**
  - dental in **tenth**
- **Coarticulation** the place of articulation for one speech sound depends on a neighbouring speech sound.
- Eg. /n/ in **ten pounds** vs. **ten kilos**
- More problematic in natural spontaneous speech

Without explicit handling, this violates the Markov assumption!

# Phonetic Context Example



"tube"




"suit"

- Via pronunciations
  - “did you” d ih jh y ah
  - “around this” ix r aw n ih s
- Via subword units Individual phone units need to deal with a lot of variability
  - Use longer units that incorporate context, eg: biphones, demisyllables, syllables
  - Use multiple models for each: context-dependent phone models
  - Context-dependent phones are termed allophones of the parent phone


# Context-dependent phone models

- **Biphones** Each phone has a unique model for its left or right context (**left biphones** and **right biphones**)
- **Triphones** Each phone has a unique model for each left and right context.
- Represent a phone  $x$  with left context  $l$  and right context  $r$  as  $l-x-r$

# Example: left biphones



# Example: triphones





# Advantages of context-dependent models

- Context-dependent models are more **specific** than context-independent models
- Increase the detail of modelling by extending the state space – but by defining multiple context dependent models, rather than more complex context independent models
- Compensate for the incorrectness of the Markov assumption
- Divide and conquer: as more context-dependent models are defined, each one becomes responsible for a smaller region of the acoustic-phonetic space
- Let the data tell us how many contexts to model

# Types of triphone models

- **Word-internal triphones** Only take account of context within words, so “don’t ask” is represented by:

sil d+oh d-oh+n oh-n+t n-t ah+s ah-s+k s-k sil


Word internal triphones result in far fewer models than cross-word models, and enable the subword sequence for a word to be known independent of the neighbouring words. But: context is not well-modelled at word boundaries.

- **Cross-word triphones** “don’t ask” is represented by:


sil sil-d+oh d-oh+n oh-n+t n-t+ah t-ah+s ah-s+k s-k+sil sil

Note that triphone context extends across words (eg unit n-t+ah)

# Recognition with context-independent units



# Context-dependent units



# Triphone models

- **How many triphones are there?** Consider a 40 phone system.  
 $40^3 = 64\,000$  possible triphones. In a cross-word system maybe 50 000 can occur
- Number of parameters:
  - 50 000 three-state HMMs, with 10 component Gaussian mixtures per state: 1.5M Gaussians
  - 39-dimension feature vectors (12 MFCCs + energy), deltas and accelerations
  - Assuming diagonal Gaussians: about 790 parameters/state
  - **Total** about 118 million parameters!
- We would need a very large amount of training data to train such a system
  - to enable robust estimation of all parameters
  - to ensure that all possible triphones are observed (more than once) in the training data

# Modelling infrequent triphones

The number of possible triphone types is much greater than the number of observed triphone tokens.

- Smoothing – combine less-specific and more-specific models
- **Parameter Sharing** – different contexts share models
  - Bottom-up – start with all possible contexts, then merge
  - Top-down – start with a single context, then split
- All approaches are data driven

NB: knowledge is used to make it work effectively

# Smoothing: Backing off

- **Basic idea** Use less-specific models when there is not enough data to train a more specific one
- For example if a triphone is not observed (or only a few examples are observed) use a biphone model:  
$$\text{sh-iy+1} \rightarrow \text{iy+1}$$
- If only a few biphone occurrences use a monophone:  
$$\text{sh-iy+1} \rightarrow \text{iy+1} \rightarrow \text{iy}$$
- Use a minimum training example count to determine whether a triphone should be modelled or backed-off to a biphone (likewise for biphones)
- Ensures that each model is well trained
- But training data is sparse (especially when cross-word triphones are used) so relatively few specific triphone models

# Smoothing: Interpolation

- **Basic idea** Combine less-specific models with more specific models
- Interpolate the parameters of a triphone  $\lambda^{tri}$  with those of a biphone  $\lambda^{bi}$  and a monophone  $\lambda^{mono}$ :

$$\hat{\lambda}^{tri} = \alpha_3 \lambda^{tri} + \alpha_2 \lambda^{bi} + \alpha_1 \lambda^{mono}$$

- Estimate the interpolation parameters  $\alpha$  using deleted interpolation
- This enables more triphone models to be estimated, but adds robustness by sharing training data from other contexts (through the biphone and monophone models)




# Parameter Sharing

- **Basic idea** Explicitly share models or parameters between different contexts
  - enables **training data** to be shared between the models
  - enables models to share parameters
- Sharing can take place at different levels
  - 1 Sharing Gaussians: all distributions share the same set of Gaussians but have different mixture weights (**tied mixtures**)
  - 2 Sharing states: allow different models to share the same states (**state clustering**)
  - 3 Sharing models: merge those context-dependent models that are the most similar (**generalised triphones**)


# Sharing Models: Generalized triphones

- **Basic idea** Merge similar context-dependent models
- Bottom-up merging: Compare allophone models with different triphone contexts and merge those that are similar
- Merged models will be estimated from more data than individual models: more accurate models, fewer models in total
- The resultant merged models are referred to as generalized triphones

# Example: Generalized Triphones




## Example : Generalized triphones




Simple triphones (no sharing)

## Example : Generalized triphones




Simple triphones (no sharing)




Generalized triphones (model sharing)

# State Clustering




Simple triphones (no sharing)

# State Clustering




Simple triphones (no sharing)




State-clustered triphones (state sharing)

# State Clustering



Simple triphones (no sharing)



State-clustered triphones (state sharing)



# Sharing States: State clustering

- **Basic idea** States which are responsible for acoustically similar data are shared
- By clustering similar states, the training data associated with individual states may be pooled together – results in better parameter estimates for the state
  - 1 Create a set of context dependent models for a parent phone
  - 2 Cluster and tie similar states, ensuring that each resultant clustered state is responsible for “enough” training data (ie setting a minimum state occupation count)
- More flexible than clustering whole models: left and right contexts may be clustered separately

- Which states should be clustered together?
- Bottom-up clustering, for triphones of the same parent phone
  - 1 Create raw triphone models for each observed triphone context
  - 2 Cluster states as before

Drawback: unable to solve **unseen triphone problem**

- Top-down clustering: start with a parent context independent model then successively split models to create context dependent models

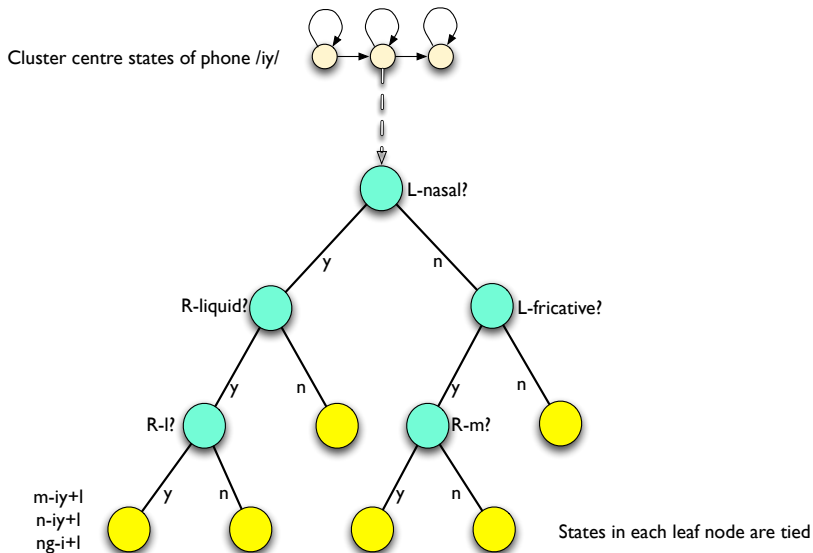
$$\text{Gain} = (L(S_1) + L(S_2)) - L(S)$$

**Phonetic decision trees**

# Phonetic Decision Trees

- **Basic idea** Build a decision tree for each state of each parent phone, with yes/no questions at each node
- At the root of the tree, all states are shared
- Questions split the pool of states, the resultant state clusters are given by the leaves of the tree
- Example questions:
  - Is the left context a nasal?
  - Is the right context a central stop?
- The questions at each node are chosen from a large set of predefined questions
- Choose the question which maximizes the likelihood of the data given the state clusters
- Stop splitting if either: (a) the likelihood does not increase by more than a predefined threshold; or (b) the amount of data associated with a split node would be below a threshold

# Phonetic Decision Tree



# Phonetic questions

- Ask questions of the form: does phone at offset  $s$  have feature  $f$ ?
- Offsets are  $+/-1$  for triphone context
- Example general questions:
  - Stop: b d g p t k
  - Nasal: m n ng
  - Fricative: ch dh f jh s sh th v z zh
  - Liquid: l r w y
  - Vowel: aa ae ah ao aw ax axr ay eh er ...
- Example consonant questions: Un/voiced, front/central/back, voiced stop, ....
- Example vowel questions: front, central, back, long, short, diphthong, rounded, ....
- Kaldi – generates all questions automatically using a top down binary clustering

# Most useful phonetic questions

- All states of all models:  
+Vowel -Vowel +Unrounded -UnFortisLenis  
+UnFortisLenis +r
- Entry state of all models:  
-UnFortisLenis -Vowel -Nasal -CentralFront  
-Unrounded -Fortis
- Exit state of all consonants:  
+Vowel +Unrounded +High +ee +Rounded +Syllabic

(for Wall St Journal read speech – Young, Odell and Woodland 1994)

# Likelihood of a state cluster (1)

- **Basic idea** Compute the log likelihood of the data associated with a pool of states
- All states pooled in a single cluster at the root
- All states have Gaussian output pdf
- Let  $S = \{s_1, s_2, \dots, s_K\}$  be a pool of  $K$  states forming a cluster, sharing a common mean  $\mu_S$  and covariance  $\Sigma_S$
- Let  $X$  be the set of training data
- Let  $\gamma_s(x)$  be the probability that  $x \in X$  was generated by state  $s$  (i.e. state occupation probability)
- The log likelihood of the data associated with cluster  $S$  is:

$$L(S) = \sum_{s \in S} \sum_{x \in X} \log P(x | \mu_S, \Sigma_S) \gamma_s(x)$$

## Likelihood of a state cluster (2)

- Don't need to iterate through all data for each state
- If the output pdfs are Gaussian it can be shown that

$$L(S) = -\frac{1}{2} \left( \log \left( (2\pi)^d |\Sigma_S| \right) + d \right) \sum_{s \in S} \sum_{x \in X} \gamma_s(x)$$

where  $d$  is the dimension of the data

- Thus  $L(S)$  depends on only
  - the pooled state variance  $\Sigma_S$  - can be computed from the means and variances of the individual states in the pool
  - and the state occupation probabilities already computed when forward-backward was carried out



# State splitting (1)

- **Basic idea** Use the likelihood of the parent state and of the split states to choose the splitting question
- Split  $S$  into two partitions  $S_y$  and  $S_n$  using a question about the phonetic context
- Each partition is now clustered together to form a single Gaussian output distribution with mean  $\mu_{S_y}$  and covariance  $\Sigma_{S_y}$  (for partition  $S_y$ )
- The likelihood of the data after partition is given by  $L(S_y) + L(S_n)$
- The total likelihood of the partitioned data will increase by

$$\Delta = L(S_y) + L(S_n) - L(S)$$

## State splitting (2)

- **Basic idea** Use the likelihood of the parent state and of the split states to choose the splitting question


$$\Delta = L(S_y) + L(S_n) - L(S)$$

- Cycle through all possible questions, compute  $\Delta$  for each and choose the question for which  $\Delta$  is biggest
- Continue by splitting each of the new clusters  $S_y$  and  $S_n$
- Terminate when
  - ① Maximum  $\Delta$  falls below a threshold
  - ② The amount of data associated with a split node falls below a threshold
- For a Gaussian output distribution: State likelihood estimates can be estimated using just the *state occupation counts* (obtained at alignment) and the parameters of the Gaussian – no need to use the acoustic data
- State occupation count: sum of state occupation probabilities for a state over time


# “Mixing up”

- **Basic idea** Transforming an HMM-based system based on Gaussian distributions to one based on mixtures of Gaussians
- The above methods for state clustering assume that the state outputs are Gaussians – this makes the computations **much** simpler
- BUT: Gaussian mixtures offer much better acoustic models than Gaussians
- Solution:
  - Perform state clustering using Gaussian distributions
  - Split the Gaussian distributions in the clustered states, by cloning and perturbing the means by a small fraction of the standard deviation, and retrain.
  - Repeat by splitting the dominant (highest state occupation count) mixture components in each state

# “Mixing up”



State-clustered triphones (Gaussians)



State-clustered triphones (GMMs)

# Summary: Context-dependent phone models

- Share parameters through state clustering
- Cluster states using phonetic decision trees for each state of parent phone
- Use Gaussian distributions when state clustering
- Then split Gaussians and retrain to obtain a GMM state clustered system

# References: context-dependent phone models

- c1980: First proposed by Bahl et al (IBM)
- Schwartz et al (1985): first paper using triphone models
- Lee (1990): generalized triphones
- Bellegarda (1990), Huang (1992): tied mixture modelling
- Bahl et al (1991): phonetic decision trees first proposed
- Young and Woodland (1994): state clustering
- Young et al (1994): decision tree-based state clustering
- Povey, 2012: Lecture on phonetic context-dependency in Kaldi